

Ayush Sharma  
(150123046)

6 April 2018

## I. Question 4.

```
CREATE TABLE twitter.table4 (  
    date date,  
    count int,  
    author_id text,  
    PRIMARY KEY (date, count, author_id)  
) WITH CLUSTERING ORDER BY (count DESC, author_id ASC)  
    AND bloom_filter_fp_chance = 0.01  
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}  
    AND comment = ''  
    AND compaction = {'class': 'org.apache.cassandra.db.compaction. -  
SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold':  
'4'}  
    AND compression = {'chunk_length_in_kb': '64', 'class': 'org.a-  
pache.cassandra.io.compress.LZ4Compressor'}  
    AND crc_check_chance = 1.0  
    AND dclocal_read_repair_chance = 0.1  
    AND default_time_to_live = 0  
    AND gc_grace_seconds = 864000  
    AND max_index_interval = 2048  
    AND memtable_flush_period_in_ms = 0  
    AND min_index_interval = 128  
    AND read_repair_chance = 0.0  
    AND speculative_retry = '99PERCENTILE';
```

## II. Question 8.

```
CREATE TABLE twitter.table8 (  
    date date,  
    count int,  
    hashtag text,  
    mention text,  
    PRIMARY KEY (date, count, hashtag, mention)  
) WITH CLUSTERING ORDER BY (count DESC, hashtag ASC, mention ASC)  
    AND bloom_filter_fp_chance = 0.01  
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}  
    AND comment = ''  
    AND compaction = {'class': 'org.apache.cassandra.db.compaction. -  
SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold':  
'4'}  
    AND compression = {'chunk_length_in_kb': '64', 'class': 'org.a-  
pache.cassandra.io.compress.LZ4Compressor'}  
    AND crc_check_chance = 1.0  
    AND dclocal_read_repair_chance = 0.1  
    AND default_time_to_live = 0  
    AND gc_grace_seconds = 864000  
    AND max_index_interval = 2048  
    AND memtable_flush_period_in_ms = 0  
    AND min_index_interval = 128  
    AND read_repair_chance = 0.0  
    AND speculative_retry = '99PERCENTILE';
```